

CS61065: Theory And Applications of Blockchain

Advanced Consensus Mechanisms

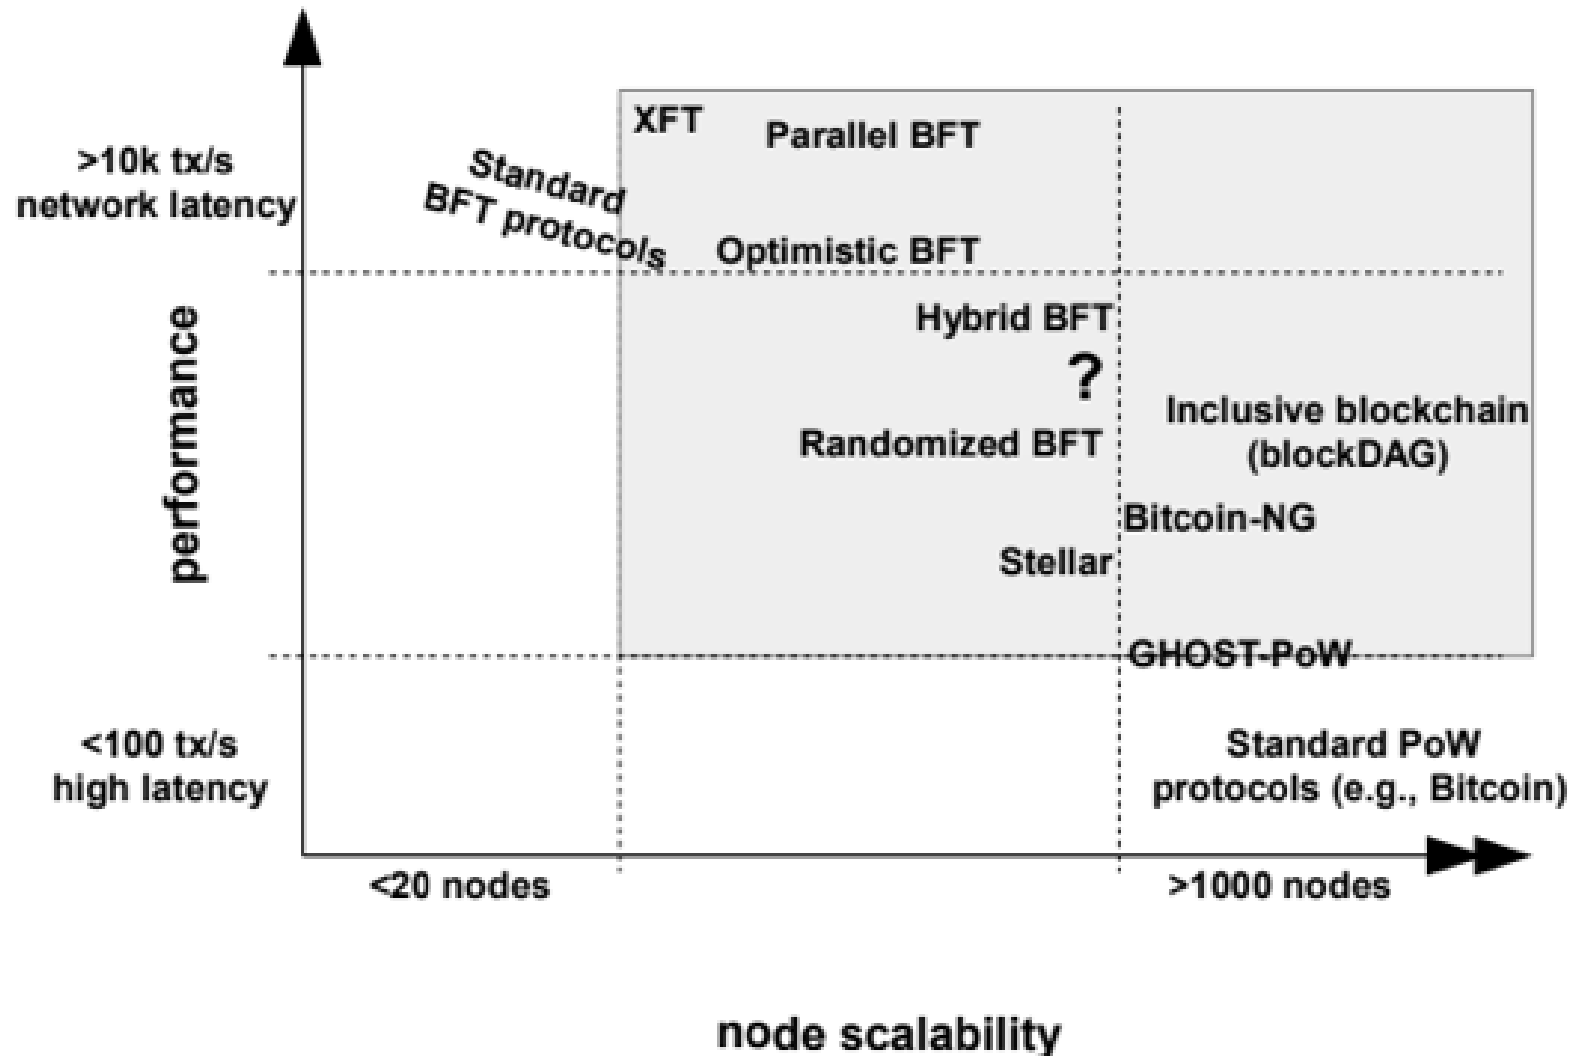
Department of Computer Science
and Engineering



INDIAN INSTITUTE OF TECHNOLOGY
KHARAGPUR

Sandip Chakraborty
sandipc@cse.iitkgp.ac.in

Consensus Scalability



Vukolić, Marko. "The quest for scalable blockchain fabric: Proof-of-work vs. BFT replication." *International Workshop on Open Problems in Network Security*. Springer, Cham, 2015.

PoW Consensus vs BFT Consensus

	PoW consensus	BFT consensus
Node identity management	open, entirely decentralized	permissioned, nodes need to know IDs of all other nodes
Consensus finality	no	yes
Scalability (no. of nodes)	excellent (thousands of nodes)	limited, not well explored (tested only up to $n \leq 20$ nodes)
Scalability (no. of clients)	excellent (thousands of clients)	excellent (thousands of clients)
Performance (throughput)	limited (due to possible of chain forks)	excellent (tens of thousands tx/sec)
Performance (latency)	high latency (due to multi-block confirmations)	excellent (matches network latency)
Power consumption	very poor (PoW wastes energy)	good
Tolerated power of an adversary	$\leq 25\%$ computing power	$\leq 33\%$ voting power
Network synchrony assumptions	physical clock timestamps (e.g., for block validity)	none for consensus safety (synchrony needed for liveness)
Correctness proofs	no	yes

Towards a Scalable Consensus

Bitcoin-NG



Eyal, I., Gencer, A. E., Sirer, E. G., & Van Renesse, R. (2016, March). **Bitcoin-NG: A Scalable Blockchain Protocol**. In *NSDI 2016*

Issues with Nakamoto Consensus

- **Transaction scalability**

- Block frequency of 10 minutes and block size of 1 MB during mining reduces the transactions supported per second


- **Issues with Forks**

- Prevents consensus finality
- Makes the system unfair - a miner with poor connectivity has always in a disadvantageous position


Bitcoin-NG: Decouple Leader Election

- Bitcoin - think of the winning miner as the **leader** - the leader serializes the transactions and include a new block in the blockchain
- Decouple Bitcoin's blockchain operations into two planes
 - **Leader election:** Use PoW to randomly select a leader (an infrequent operation)
 - **Transaction Serialization:** The leader serializes the transaction until a new leader is elected

Bitcoin vs Bitcoin-NG



Bitcoin-NG



Bitcoin-NG : Key Blocks

- Key blocks are used to choose a leader (similar to Bitcoin)
- A key block contains
 - The reference to the previous block
 - The current Unix time
 - A coinbase transaction to pay of the reward
 - A target hash value
 - A nonce field

Key Blocks

- Key blocks are generated based on regular Bitcoin mining procedure
 - Find out the nonce such that the block hash is less than the target value
- Key blocks are generated infrequently - the intervals between two key blocks is exponentially distributed

Bitcoin-NG



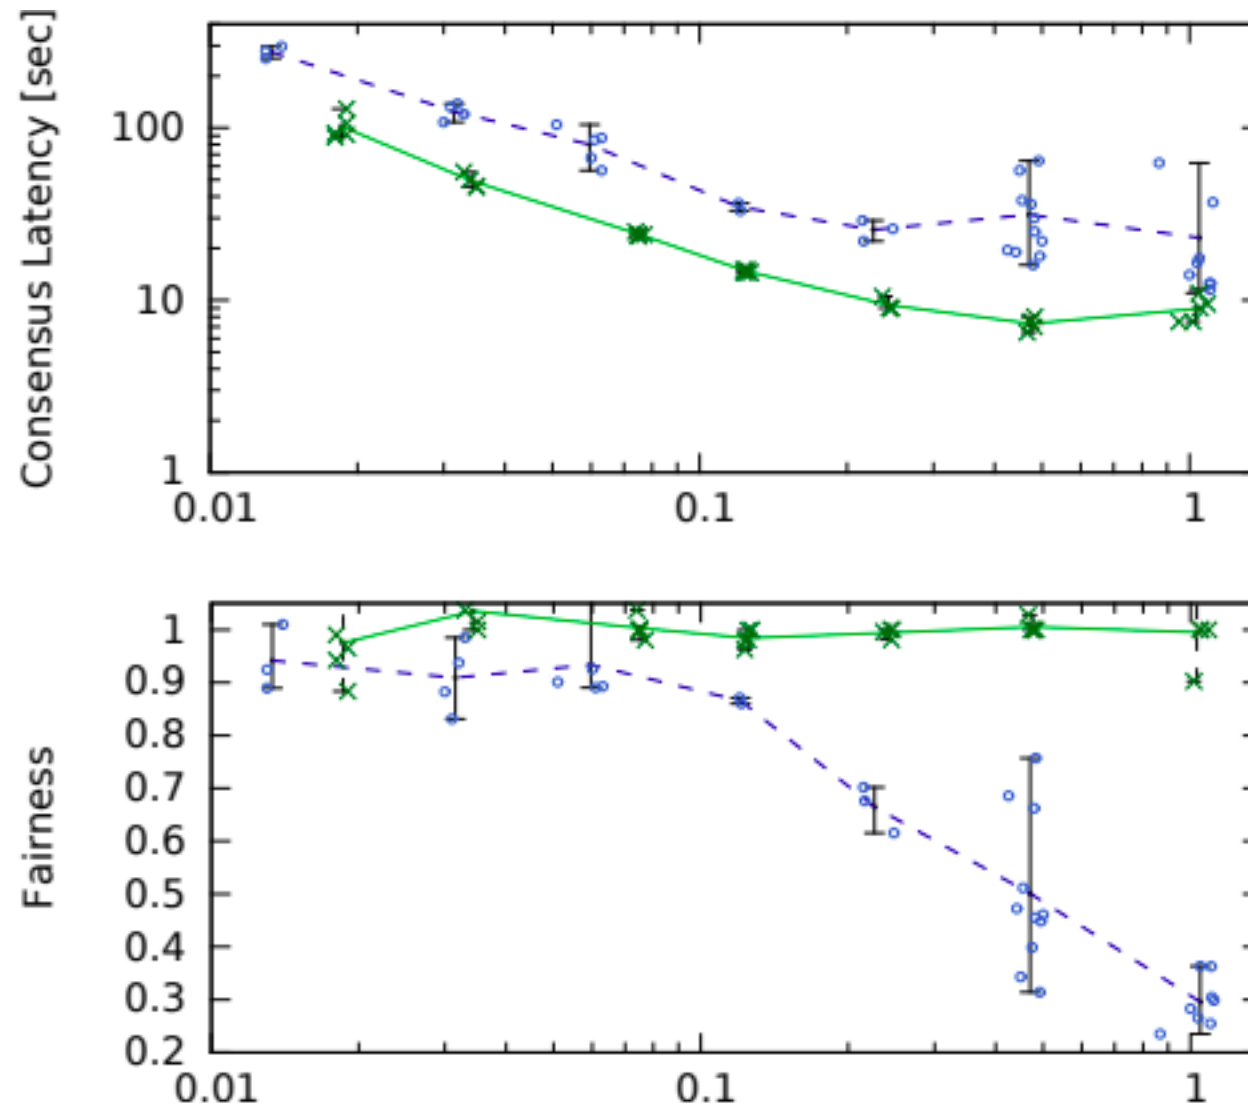
Bitcoin-NG : Microblocks

- Once a node generates a key block, it becomes the **leader** and generates further microblocks
 - Microblocks are generated at a set rate smaller than a predefined maximum
 - The rate is much higher than the key block generation rate
- A microblock contains
 - Ledger entries
 - Header
 - Reference to the previous block
 - The current Unix time
 - A cryptographic hash of the ledger entries (Merkle root)
 - A cryptographic signature of the header (signature of the key block miner)

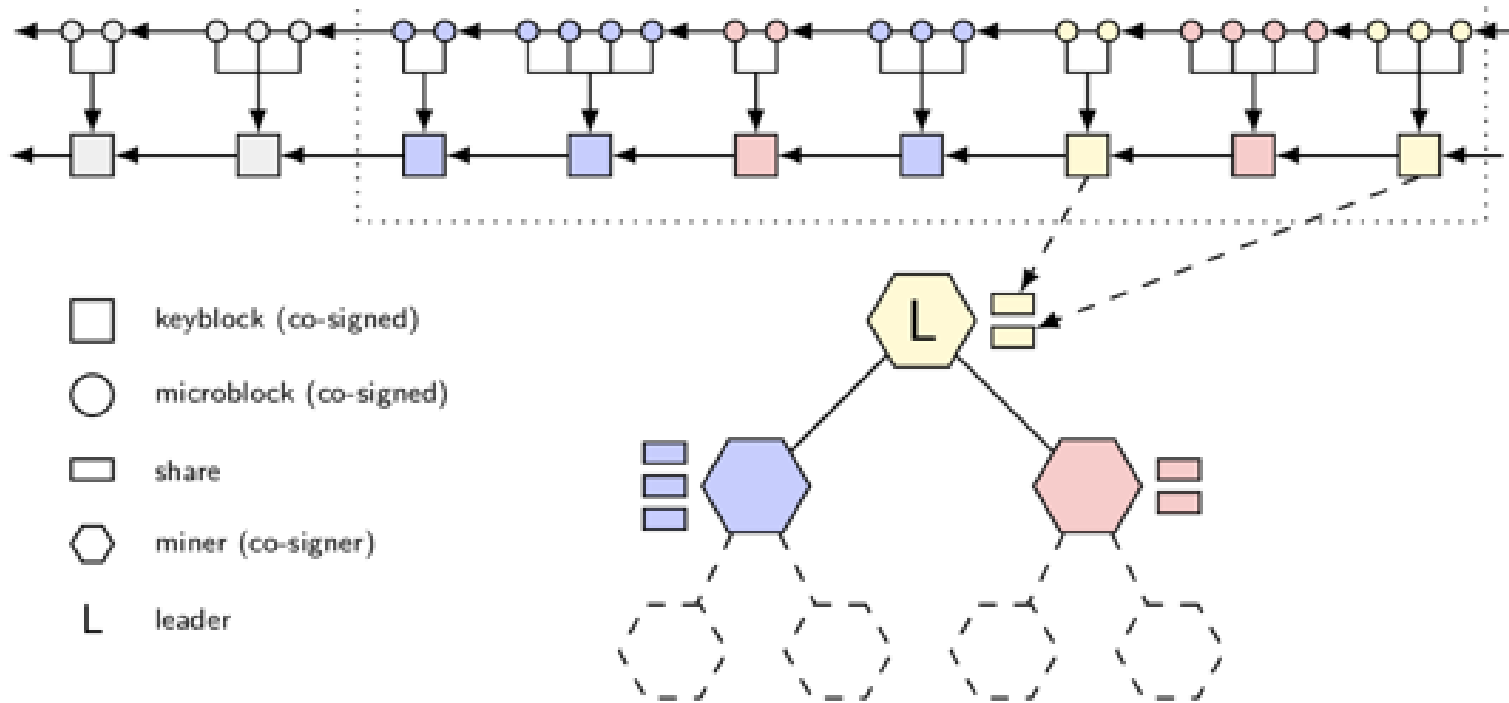
Microblock Fork

- When a miner generates a key block, he may not have heard of all microblocks generated by the previous leader
 - Common if microblock generation is frequent
 - May result in microblock fork
- A node may hear a forked microblock but not the new key block
 - This can be prevented by ensuring the reception of the key block
 - When a node sees a microblock, it waits for propagation time of the network, to make sure it is not pruned by a new key block

Bitcoin-NG Performance



Combining PoW with PBFT : Byzcoin



Kogias, E. K., Jovanovic, P., Gailly, N., Khoffi, I., Gasser, L., & Ford, B. (2016, August). **Enhancing bitcoin security and performance with strong consistency via collective signing.** In *25th USENIX Security Symposium 2016*

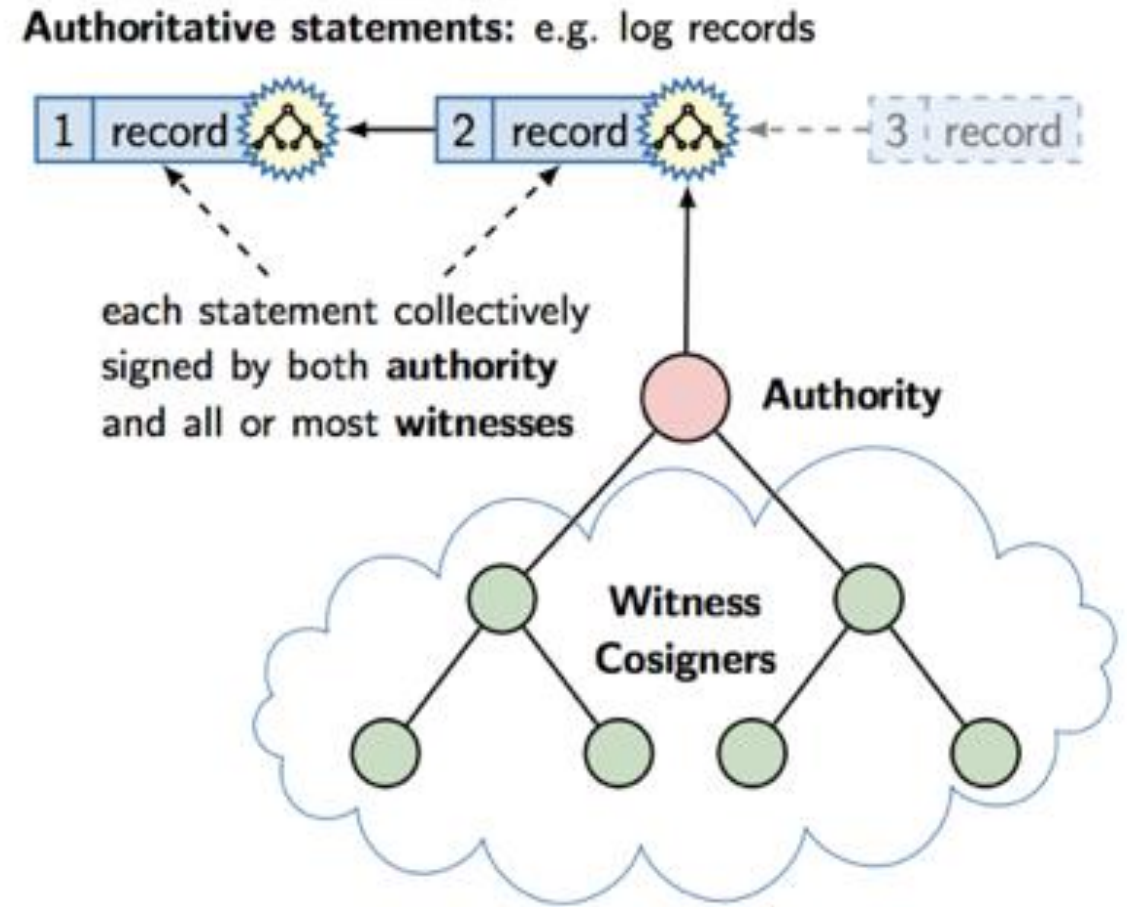
Collective Signing (CoSi)

- Method to protect “authorities and their clients” from undetected misuse or exploits
- A **scalable witness cosigning protocol** ensuring that every authoritative statement is validated and publicly logged by a diverse group of witnesses before any client accepts it
- A statement S collectively signed by W witnesses assures clients that S has been seen, and not immediately found erroneous, by those W observers.

Syta, Ewa, *et al.* "Keeping authorities "honest or bust" with decentralized witness cosigning" *2016 IEEE Symposium on Security and Privacy (SP)*, 2016.

CoSi Architecture

- The leader organizes the witnesses in a tree structure – a scalable way of aggregating signatures coming from the children
- Three rounds of PBFT (pre-prepare, prepare and commit) can be simulated using two rounds of CoSi protocol



Schnorr Multi-signature

- The basic CoSi protocol uses **Schnorr multisignatures**, that rely on a group G of prime order
 - *Discrete logarithmic problem is believed to be hard*
- **Key Generation:**
 - Let G be a group of prime order r . Let g be a generator of G .
 - Select a random integer x in the interval $[0, r - 1]$. x is the private key and g^x is the public key.
 - N signers with individual private keys x_1, x_2, \dots, x_N , and the corresponding public keys $g^{x_1}, g^{x_2}, \dots, g^{x_N}$

Schnorr Multi-signature

- **Signing:**
 - Each signer picks up the random secret v_i , generates $V_i = g^{v_i}$
 - The leader collects all such V_i , aggregates them $V = \prod V_i$, and uses a hash function to compute a collective challenge $c = H(V || S)$. This challenge is forwarded to all the signers.
 - The signers send the response $r_i = v_i - cx_i$. The leader computes the aggregated as $r = \sum r_i$. The signature is (c, r) .

Schnorr Multi-signature

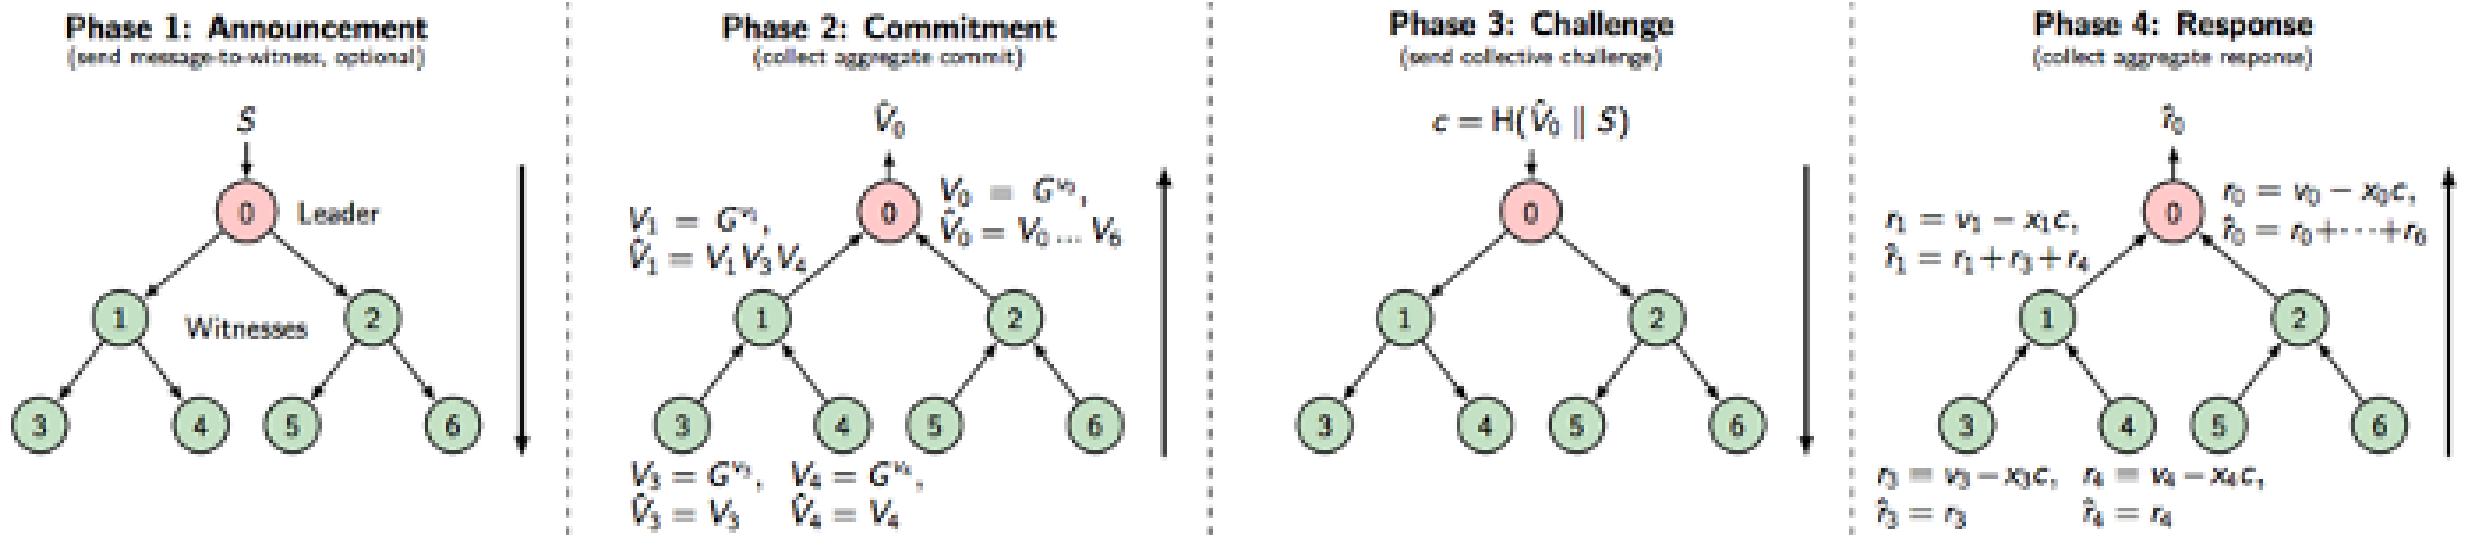
- **Verification:**
 - The verification key is $y = \prod g^{x_i}$
 - The signature is (c, r) , where $c = H(V||S)$ and $r = \sum r_i$
 - Let $V_v = g^r y^c$
 - Let $r_v = H(V_v||S)$
 - If $r_v = r$, then the signature is verified

Schnorr Multi-signature

- **Proof:**

- The verification key is $y = \prod g^{x_i}$
- The signature is (c, r) , where $c = H(V || S)$ and $r = \sum r_i$
- $V_v = g^r y^c = g^{\sum (v_i - cx_i)} \prod g^{cx_i} = g^{\sum (v_i - cx_i)} g^{\sum cx_i} = g^{\sum v_i} = \prod g^{v_i} = \prod V_i = V$
- So, $r_v = H(V_v || S) = H(V || S) = r$

CoSi Protocol



- One CoSi round to implement PBFT's pre-prepare and prepare phases
- Second CoSi round to implement PBFT's commit phase

Revisiting the Requirements for Blockchain Consensus

- **Byzantine fault tolerant** – the system should work even in the presence of malicious users while operating across multiple administrative domains
- Should provide **strong consistency guarantee** across replicas
- Should **scale well to increasing workloads** in terms of transactions processed per unit time
- Should **scale well to increasing network size**

Bitcoin-NG: The Issue with a Faulty Key Block

- **Problem with Bitcoin-NG:** A faulty key block is verified only after end of the round
 - A faulty miner can introduce a number of correct microblocks following a faulty microblock in the system - certainly a overhead for the application - **a fork alleviates the problem further**

Bitcoin-NG: The Issue with a Faulty Key Block

- **Problem with Bitcoin-NG:** A faulty key block is verified only after end of the round
 - A faulty miner can introduce a number of correct microblocks following a faulty microblock in the system - certainly a overhead for the application - **a fork alleviates the problem further**

Solve this problem by a set of **PBFT verifiers - who will verify a block and then only the block is added in the Blockchain**

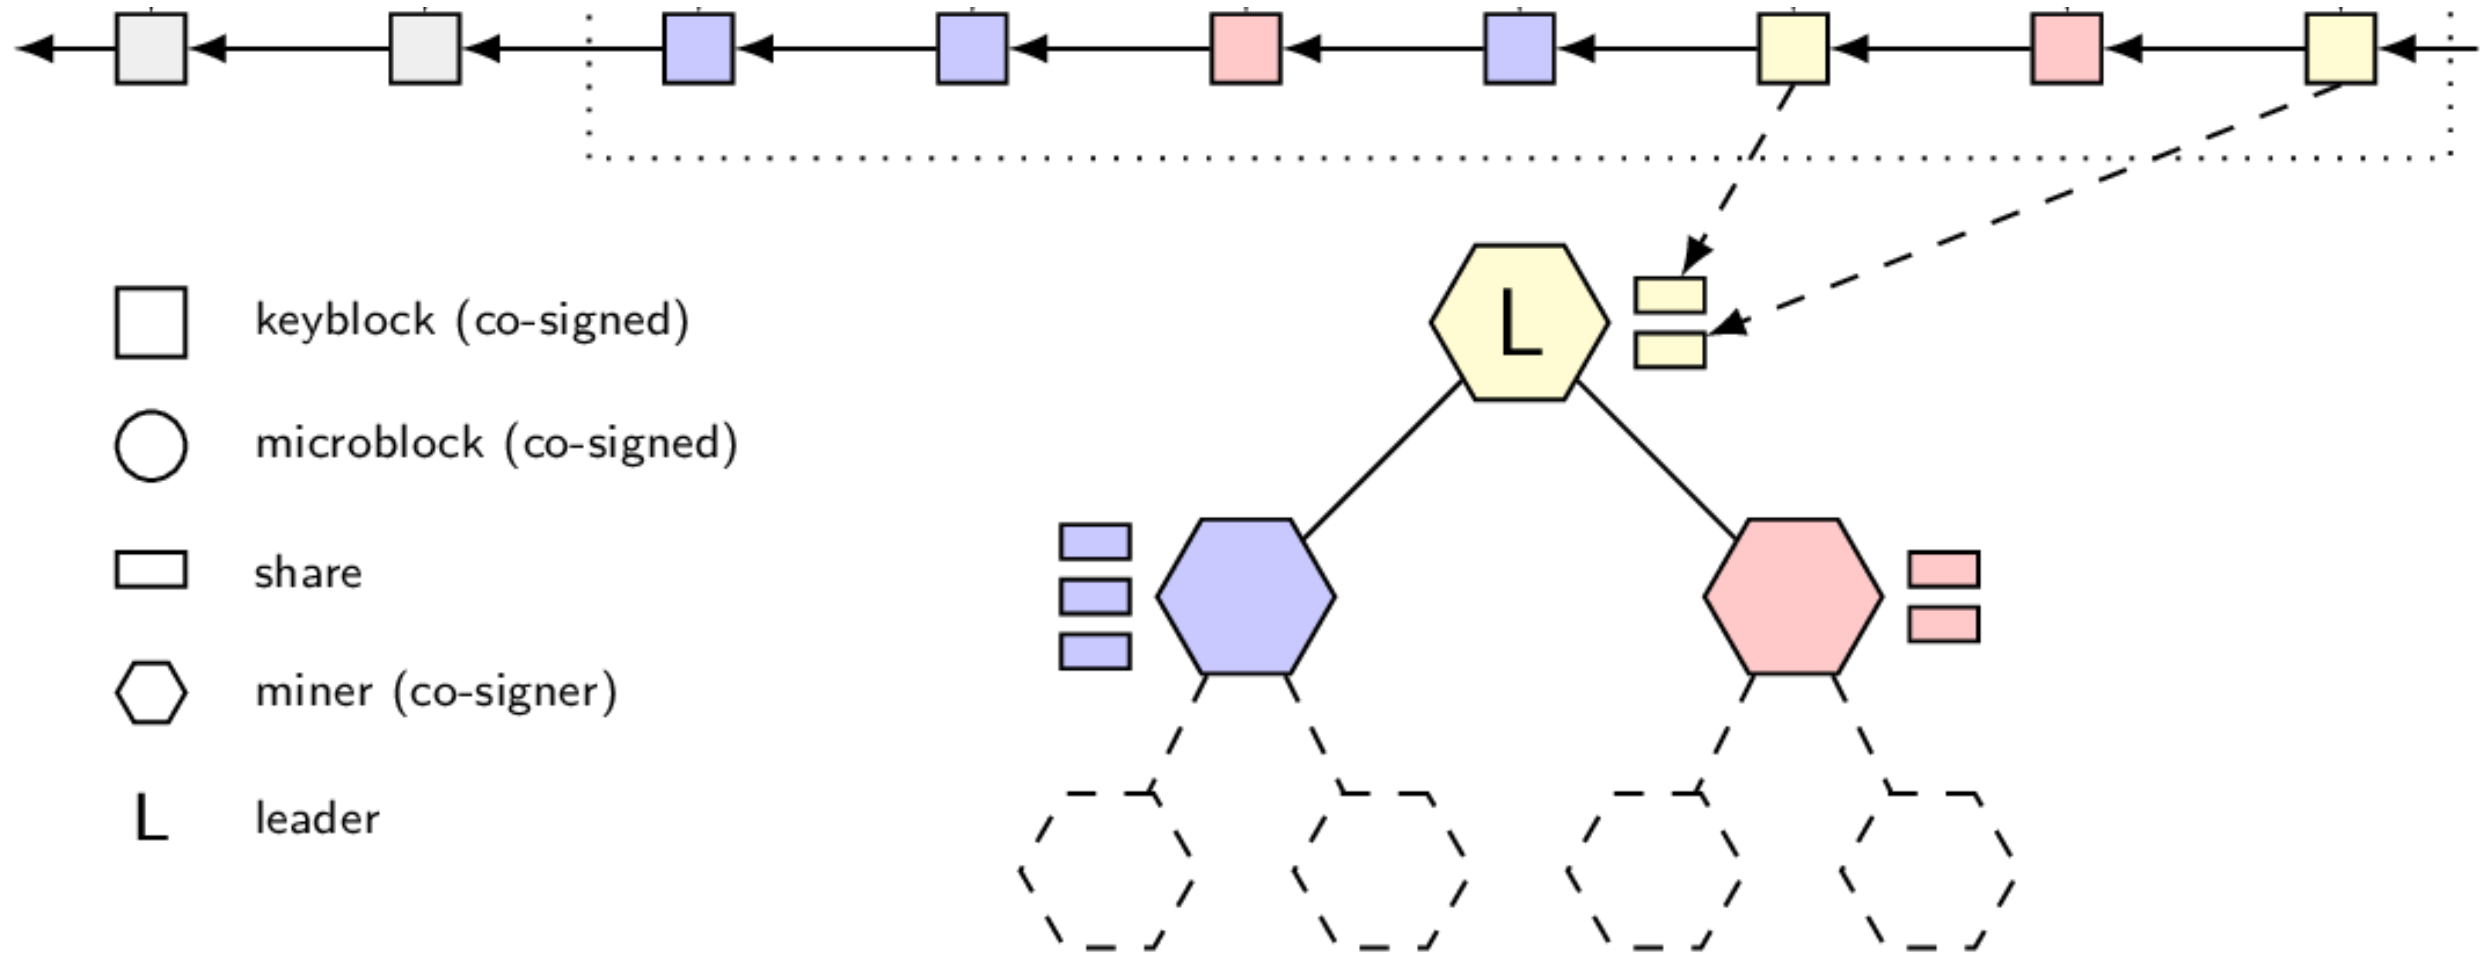
Issues with PBFT

- PBFT requires a **static consensus group** (because of message passing)
- **Scalability** (in terms of nodes) is a problem for PBFT
 - $O(n^2)$ communication complexity
 - $O(n)$ verification complexity
 - Absence of third-party verifiable proofs (PBFT uses MAC - need to share the keys among the miners)
- **Sybil attack** - create multiple pseudonymous identities to subvert the **$3f+1$** requirements of PBFT

Open the Consensus Group

- Use PoW based system to give a *proof of membership* of a miner as a part of the trustees
- Maintains a “balance of power” within the BFT consensus group
 - Use a fixed-size sliding window
 - Each time a miner finds a new block, it receives a *consensus group share*
 - The share proves the miner’s membership in the trustee group

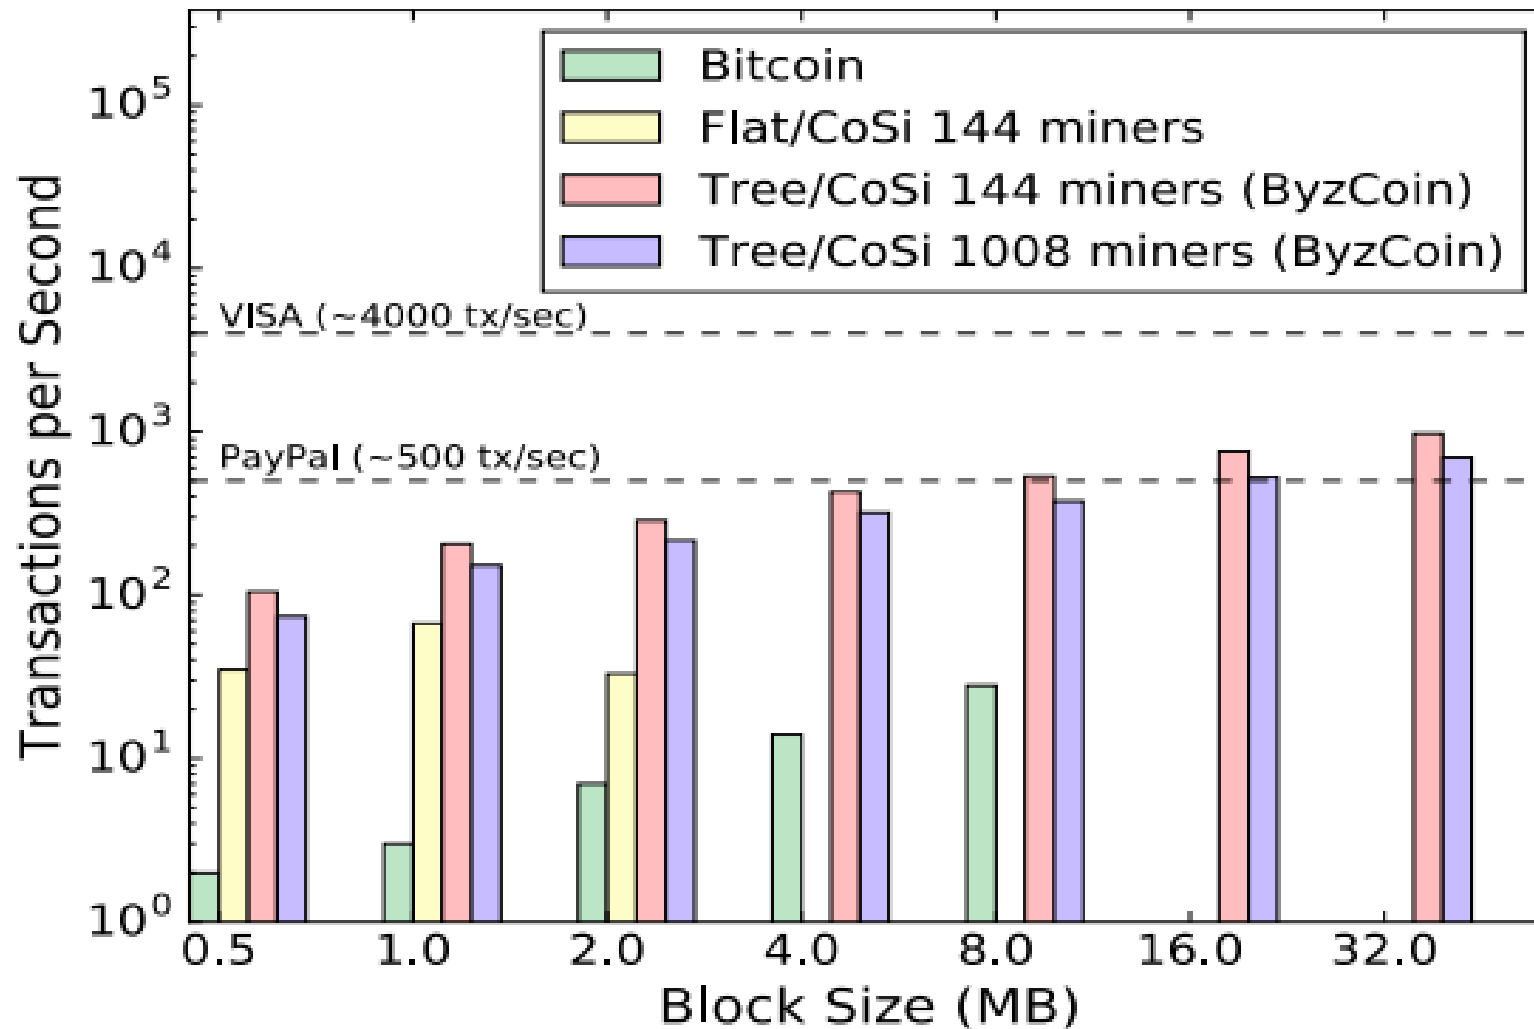
Merging BFT Consensus with PoW



Improving the Efficiency of BFT Consensus

- **Improve $O(n)$ communication complexity**
 - Use tree-based multicast protocol - share information with $O(\log n)$
- **Improve $O(n)$ complexity for verification**
 - Use Schnorr multi-signatures
 - Verification can be done in $O(1)$ through signature aggregation
- Multi-signatures + Communication trees = **CoSi**

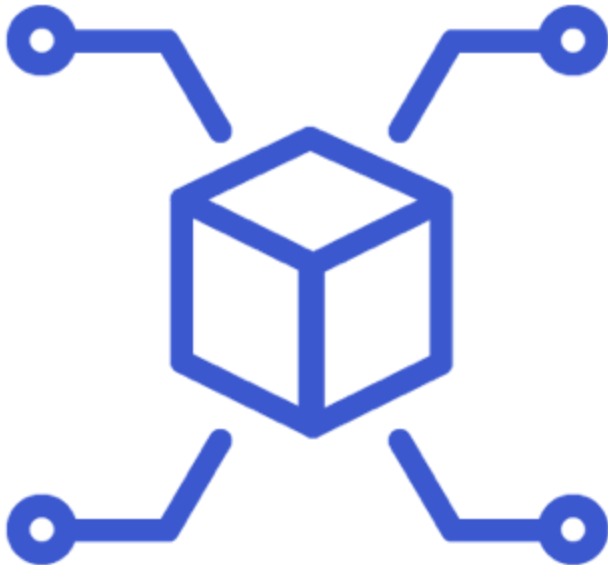
Byzcoin Performance



Algorand



Gilad, Y., Hemo, R., Micali, S., Vlachos, G., & Zeldovich, N. (2017, October). **Algorand: Scaling byzantine agreements for cryptocurrencies**. In *Proceedings of the 26th Symposium on Operating Systems Principles* (pp. 51-68). ACM.



Algorand: Scaling Byzantine Agreements for Cryptocurrencies

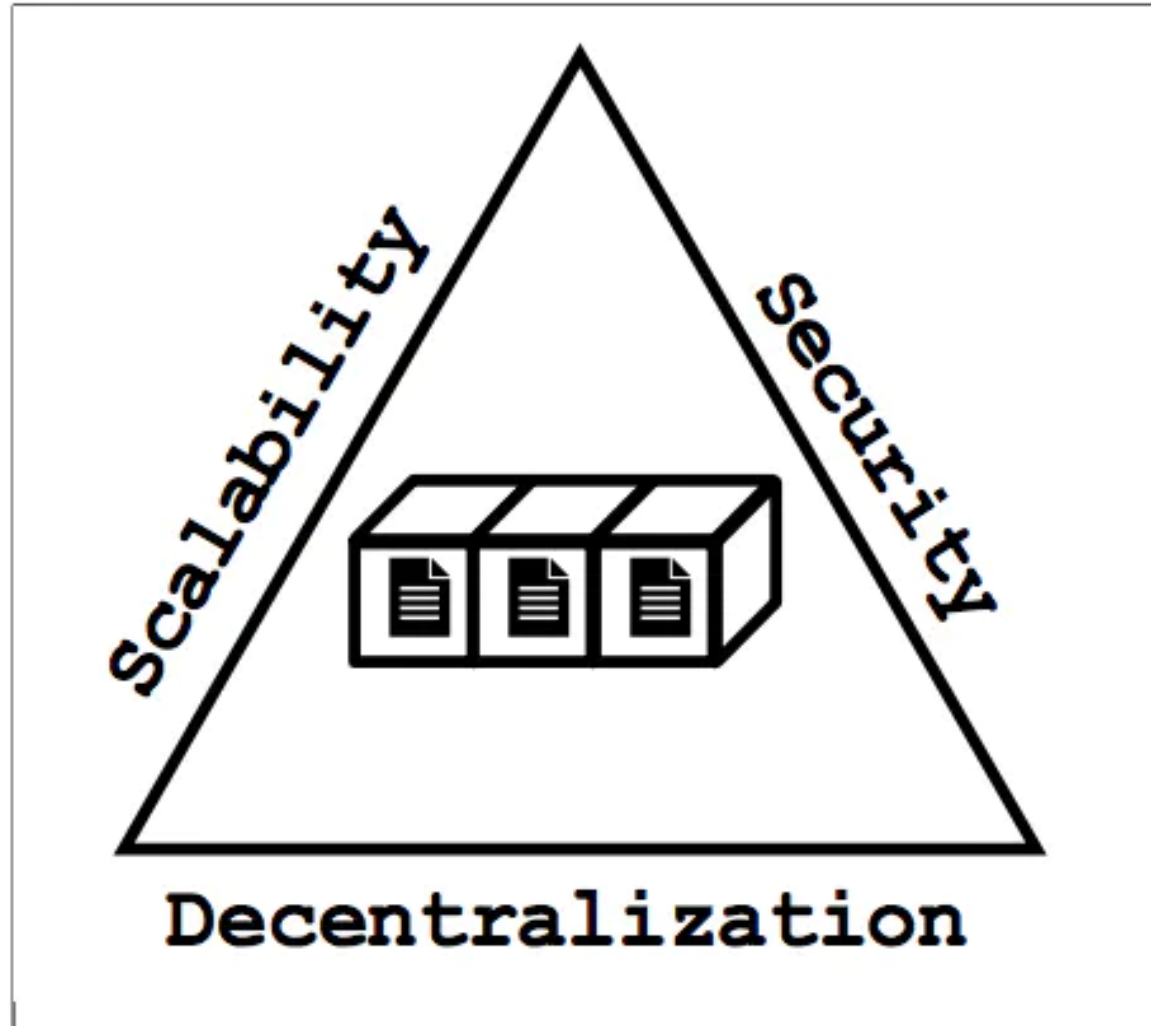
Bitcoin: Recap

- **Key Idea:**
 - Consensus through proof-of-work (PoW)
- **Communication:**
 - Gossip protocol
- **Key Assumption:**
 - Honest majority of mining computation power

Bitcoin: Limitations

- **Resource wastage:**
 - high computational, electricity cost
- **Concentration of power**
 - only ~5 mining pools control the entire system
- **Vulnerable**
 - easy to track miners, concentrated to a few mining pools - <https://www.blockchain.com/btc/blocks?page=1>
- **Scalability**
 - number of users not clear (1M, 10M, 100M??), high latency(~10minutes)
- **Ambiguity**
 - fork in blockchain

The Blockchain Performance Triangle



Algorand: Overview

- **Key Idea:**
 - Consensus through Byzantine Agreement Protocol
- **Communication:**
 - Gossip protocol
- **Key Assumption:**
 - Honest majority of money

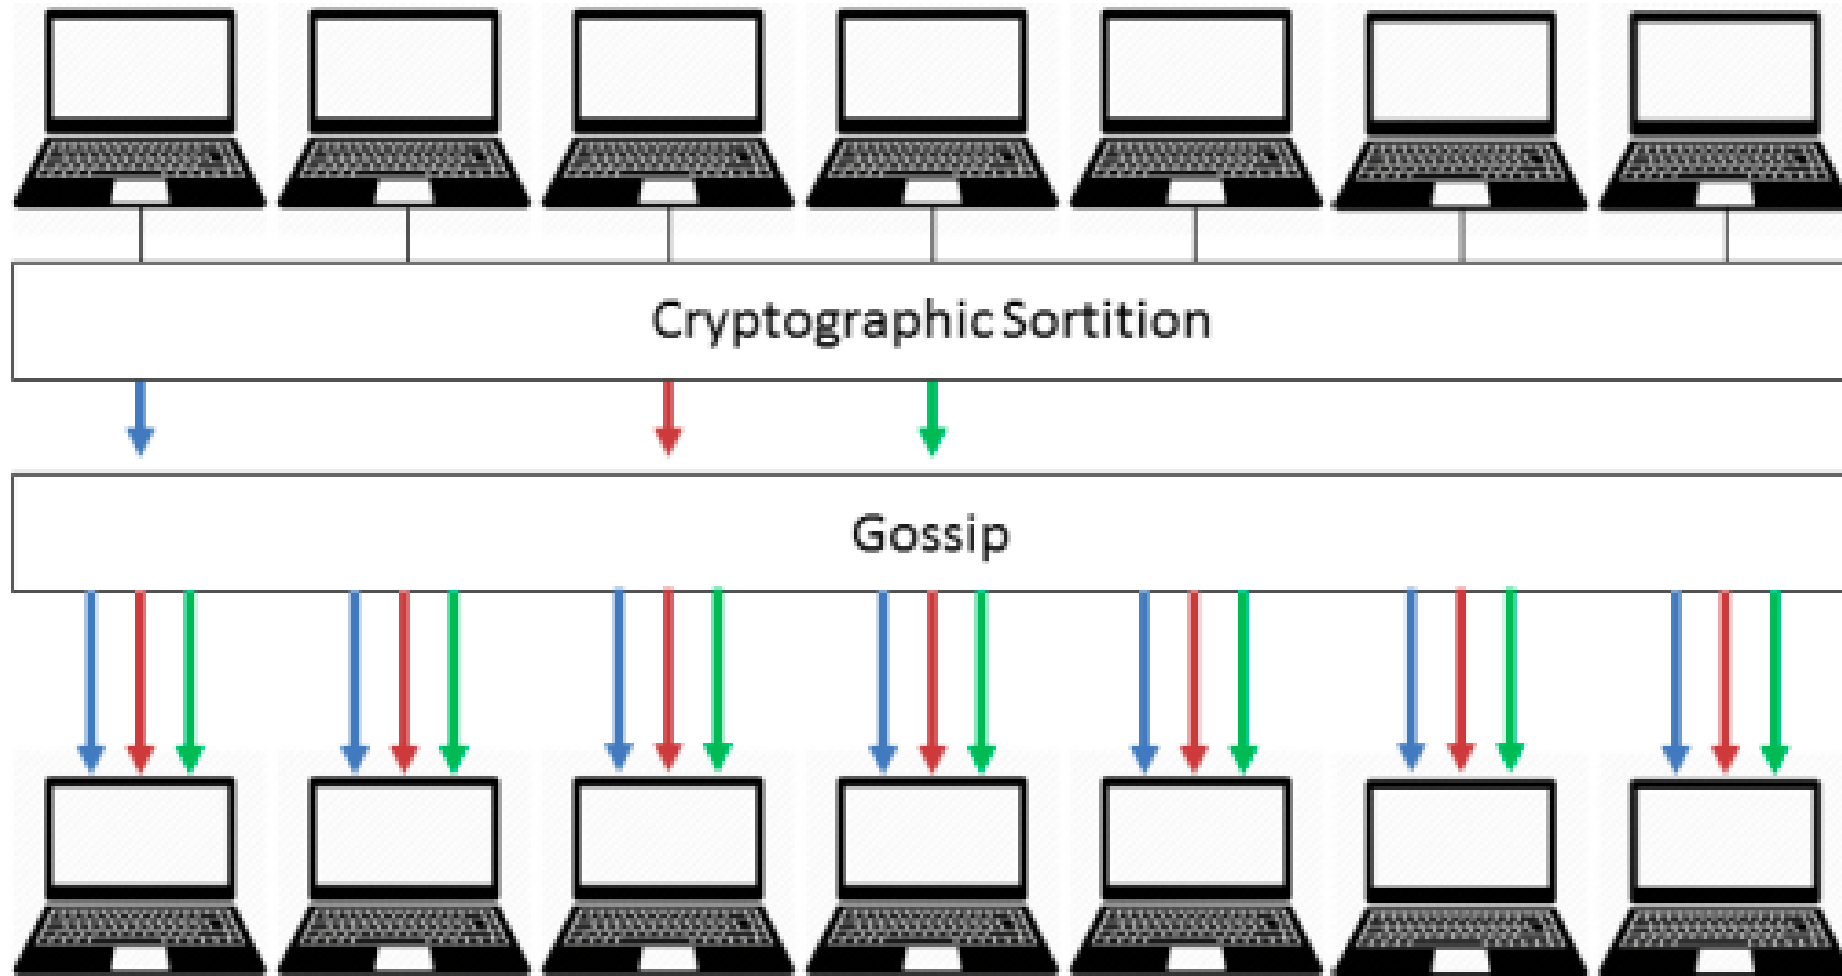
Algorand: Technical Advancement

- **Trivial computation**
 - simple operation like add, count
- **True decentralization**
 - no concentration of mining pool power, all equal miners and users
- **Finality of payment**
 - fork with very low probability, block appears and payment fixed forever
- **Scalability**
 - millions of users, only network latency (~1minute)
- **Security**
 - against bad adversary

Architecture of Algorand

- Select a random user
 - prepare a block
 - propagate block through gossiping
- Select random committee with small number of users (~10k)
 - run Byzantine Agreement on the block
 - digitally sign the result
 - propagate digital signatures
- **Who select the committee?**

Cryptographic Sortition in Algorand



Cryptographic Sortition

- Each committee member selects himself according to per-user weights
 - Implemented using Verifiable Random Functions (VRFs)
- $\langle \text{hash}, \text{proof} \rangle \leftarrow \text{VRF}_{\text{sk}}(x)$
 - **x**: input string
 - **(pk_i, sk_i)**: public/private key pair
 - **hash**: hashlenbit-long value that is uniquely determined by sk and x
 - **proof**: enables to check the hash indeed corresponds to x

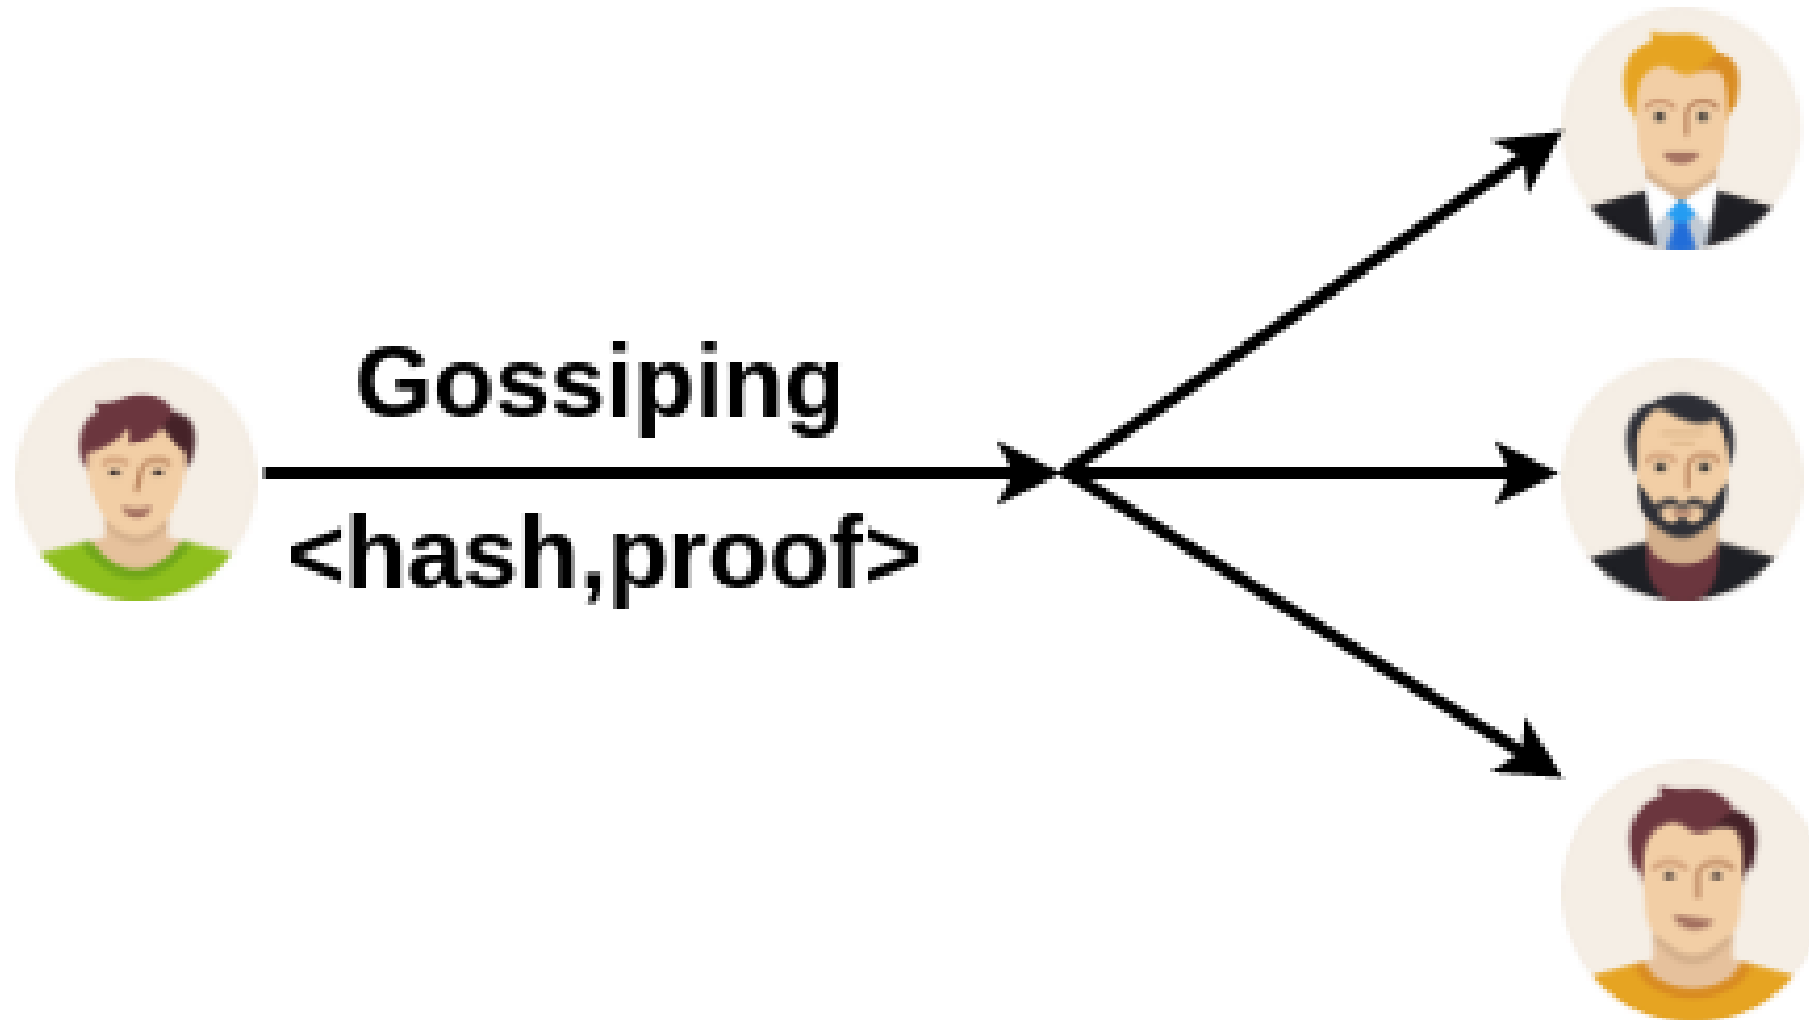
Cryptographic Sorition in Algorand: Committee Member Selection



$\langle \text{hash}, \text{proof}, j \rangle \leftarrow \text{Sortition}(\text{sk}, \text{seed}, \text{threshold}, \text{role}, w, W)$

- **seed:** publicly known random value
 - seed published at Algorand's round r using VRFs with the seed of the previous round $r - 1$
- **threshold:** determines the expected number of users selected for that role
- **role:** user for proposing a block/ committee member
- **w:** weight of a user
- **W:** weight of all users
- **j:** user gets to participate as j different "sub-users."

Cryptographic Sortition in Algorand



Cryptographic Sortition: Proof Verification



- Check if that user was selected
 - number of selected sub-users
 - zero if the user was not selected at all

Byzantine Agreement in Algorand – BA*

- **Two phase:**
 - Two phase agreement – Final Consensus and Tentative Consensus
- **Strong Synchrony:** Most honest users (say, 95%) can send message that will be received by most other honest users within a known time bound
 - Adversary can not control the network for long
 - Ensures liveness of the protocol
- **Weak Synchrony:** The network can be asynchronous for long (entirely controlled by adversary) but bounded period of time
 - There must be a strong synchrony period after a weak synchrony period
 - Algorand is safe under weak synchrony

Byzantine Agreement in Algorand – BA*

- **Two phase:**
 - Two phase agreement – Final Consensus and Tentative Consensus
- **Strong Synchrony:** Most honest users (say, 95%) can send message that will be received by most other honest users within a known time bound
 - Adversary can not control the network for long
 - Ensures liveness of the protocol
- **Weak Synchrony:** The network can be asynchronous for long (entirely controlled by adversary) but bounded period of time
 - There must be a strong synchrony period after a weak synchrony period
 - Algorand is safe under weak synchrony

Byzantine Agreement in Algorand – BA*

- **Two phase:**
 - Two phase agreement – Final Consensus and Tentative Consensus
- **Strong Synchrony:** Most honest users (say, 95%) can send message that will be received by most other honest users within a known time bound
 - Adversary can not control the network for long
 - Ensures liveness of the protocol
- **Weak Synchrony:** The network can be asynchronous for long (entirely controlled by adversary) but bounded period of time
 - **There must be a strong synchrony period after a weak synchrony period**
 - Algorand is **safe** under weak synchrony

Final Consensus

- One user reaches final consensus
 - Any other user that reaches final or tentative consensus in the same round must agree on the same block value (ensures safety)
 - Confirm a transaction when the block reaches to the final consensus



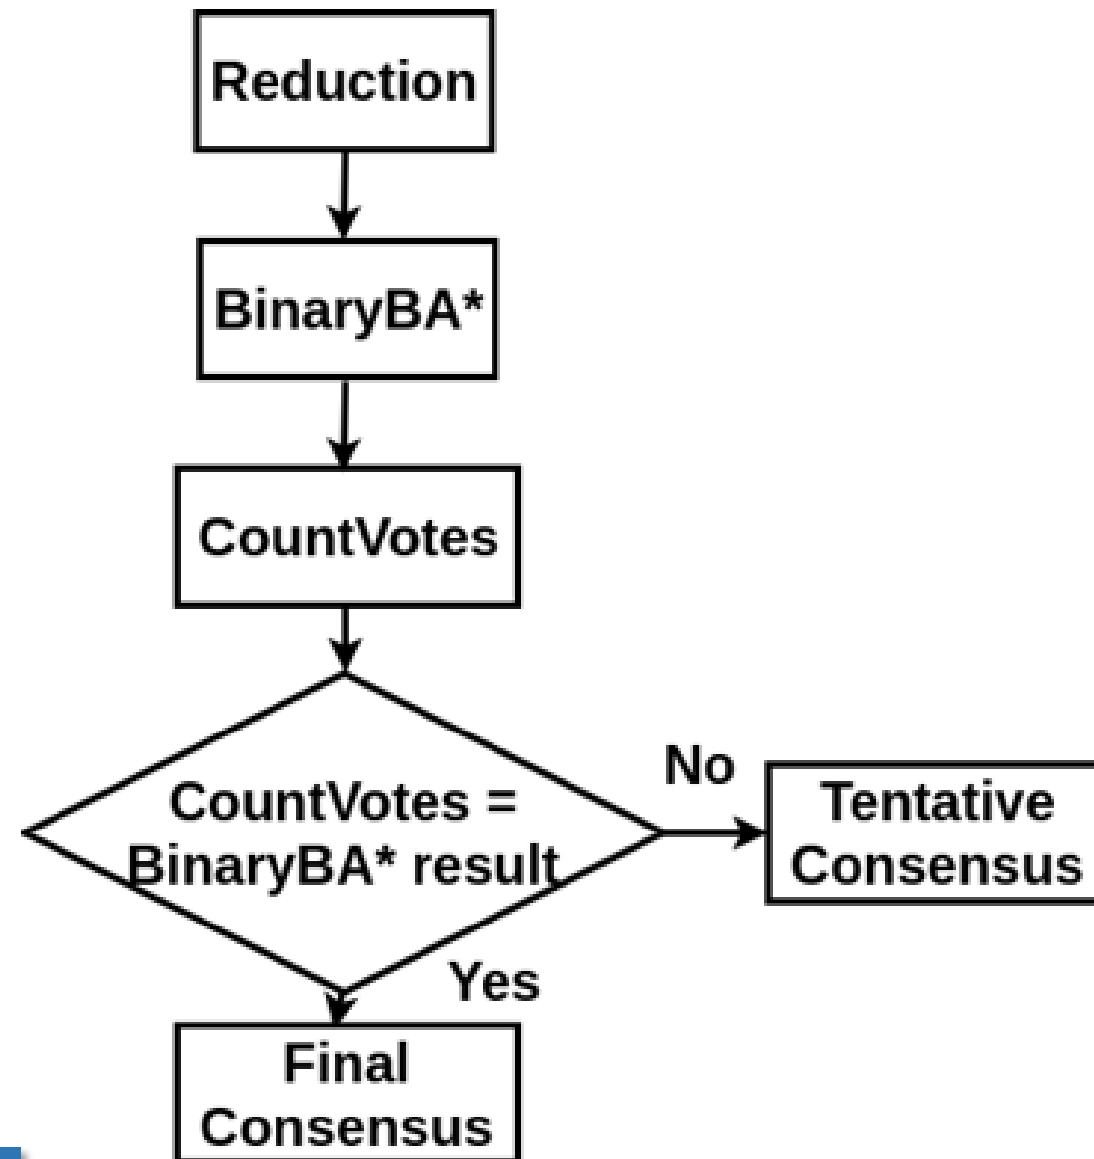
Tentative Consensus

- One user reaches tentative consensus
 - Other users may have reached consensus on a **different (but correct)** block
 - Can be in two cases
 - **The network is strongly synchronous** - adversary may be able to cause BA^* to reach tentative consensus on a block - BA^* is unable to confirm that the network was strongly synchronous
 - **The network was weakly synchronous** - BA^* can form multiple forks and reach tentative consensus on two different blocks - users are split into groups

Coming Out of Tentative Consensus

- Run BA* periodically to come out of tentative consensus - run the next round
 - Network can not be under weak synchrony all the times
 - Cryptographic sortition ensures different committee members at different rounds of the BA*

BA* Overview



thank you!